

Computer Networks

COL 334/672

Transport Layer


Tarun Mangla

Slides adapted from KR

Sem 1, 2024-25

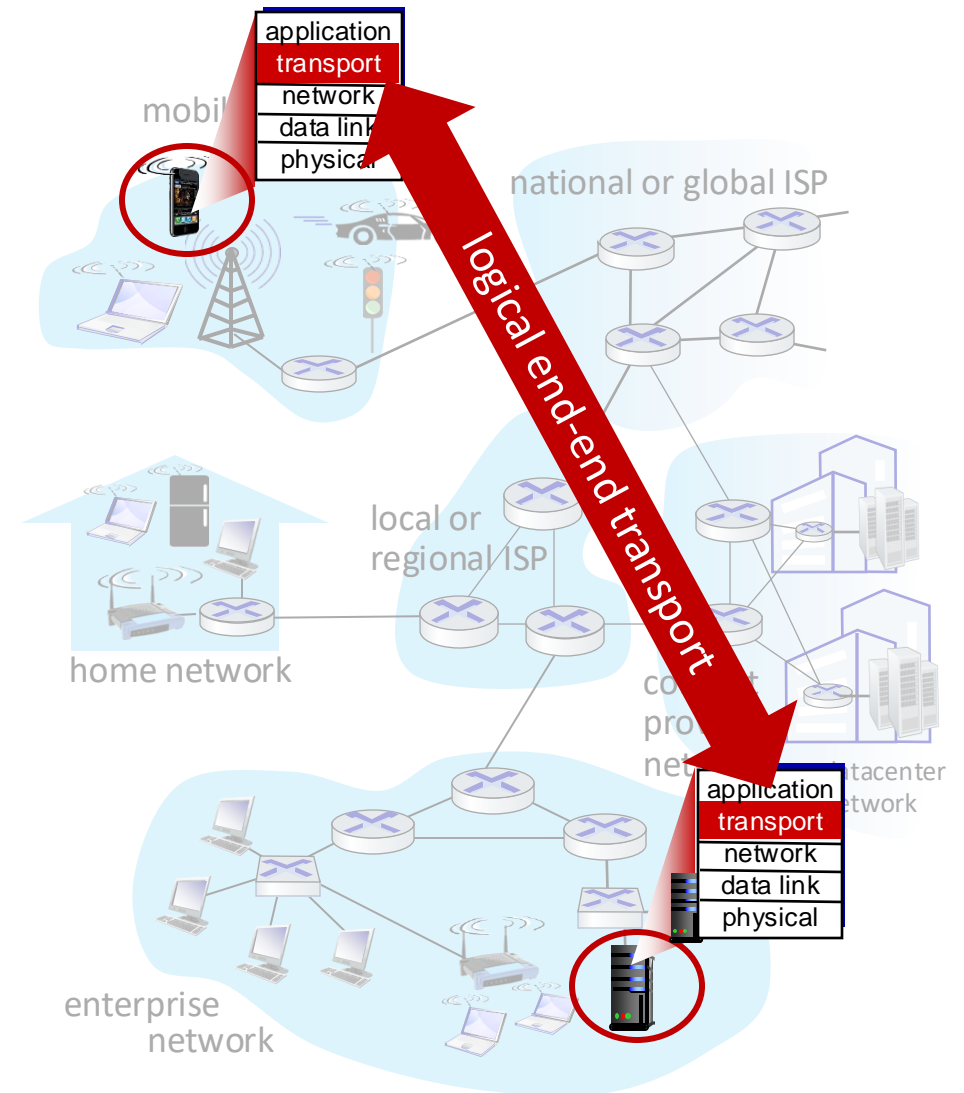
Course Recap

- *application*: supporting network applications
 - HTTP, IMAP, SMTP, DNS
- *transport*: process-process data transfer
 - TCP, UDP
- *network*: routing of datagrams from source to destination
 - IP, routing protocols
- *link*: data transfer between neighboring network elements
 - Ethernet, 802.11 (WiFi), PPP
- *physical*: bits “on the wire”



Transport services and protocols


- provide *logical communication* between application processes running on different hosts
- transport protocols actions in end systems:
 - sender: breaks application messages into *segments*, passes to network layer
 - receiver: reassembles segments into messages, passes to application layer
- two transport protocols available to Internet applications
 - TCP, UDP



Transport Layer Actions

Sender:

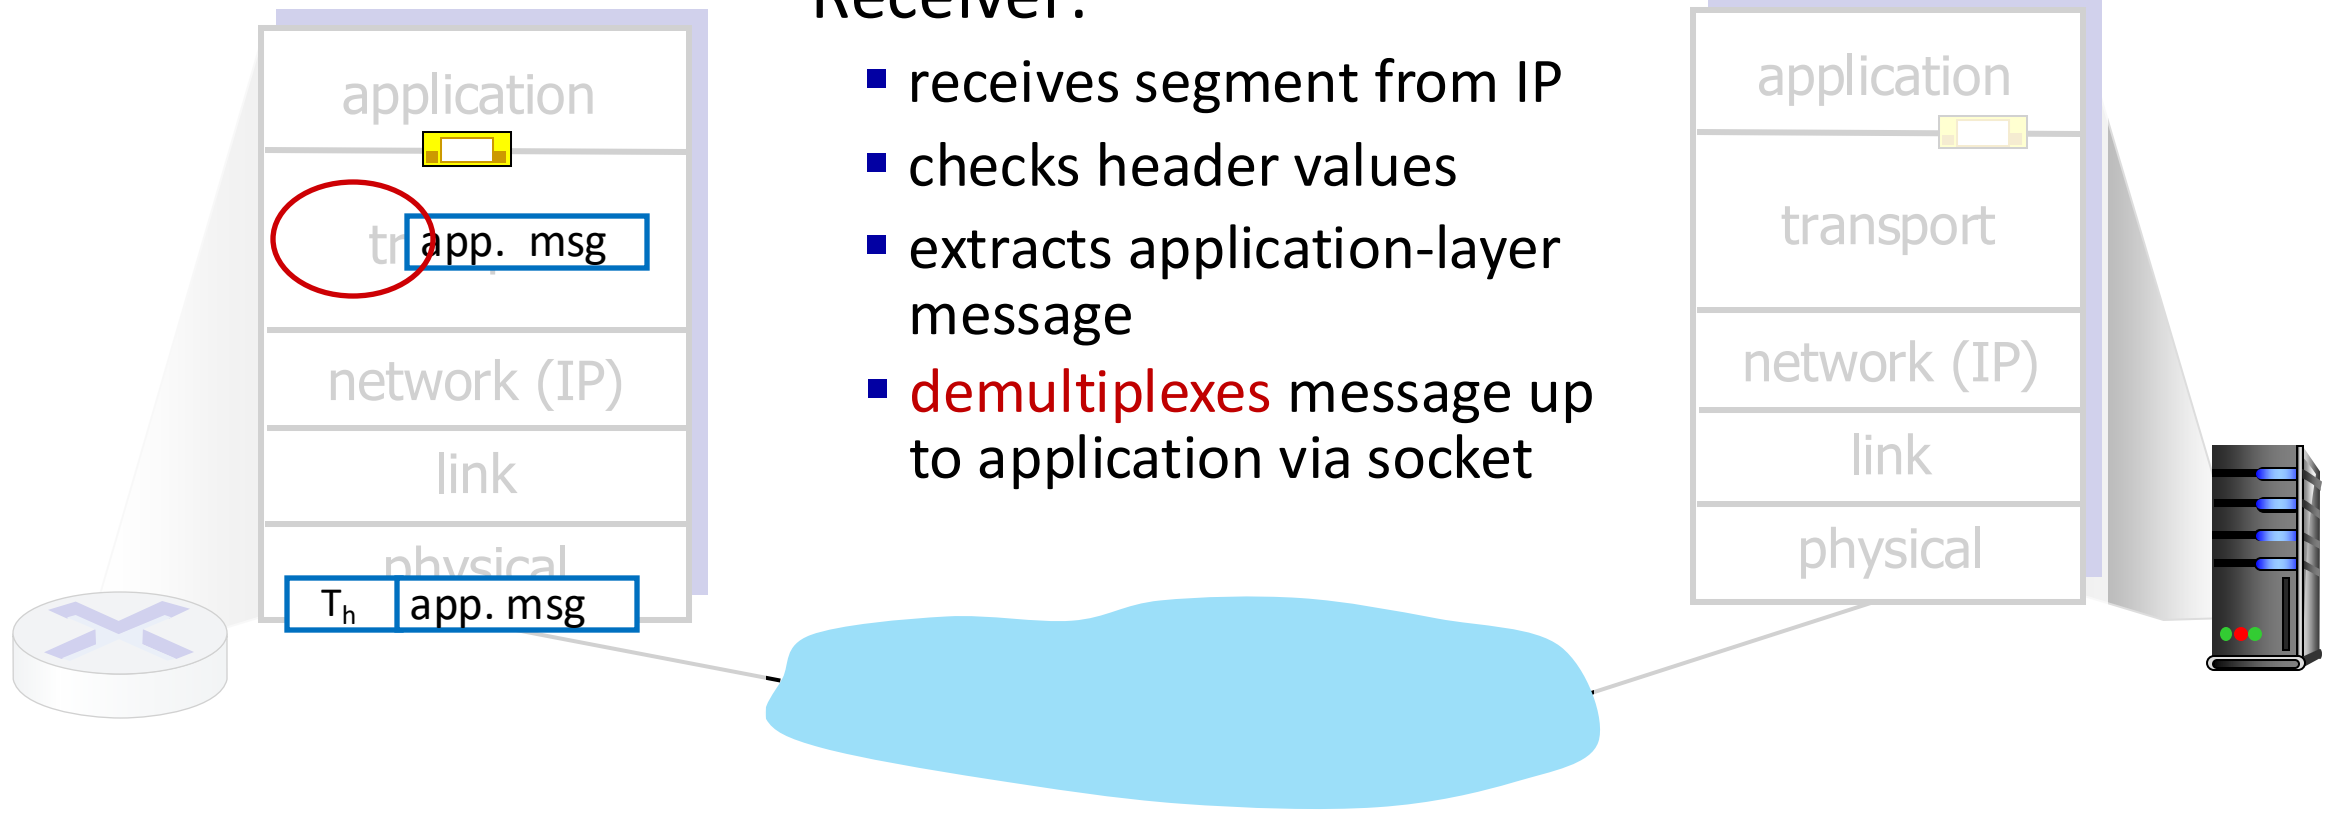
- is passed an application-layer message
- determines segment header fields values
- creates segment
- passes segment to IP



Transport Layer Actions

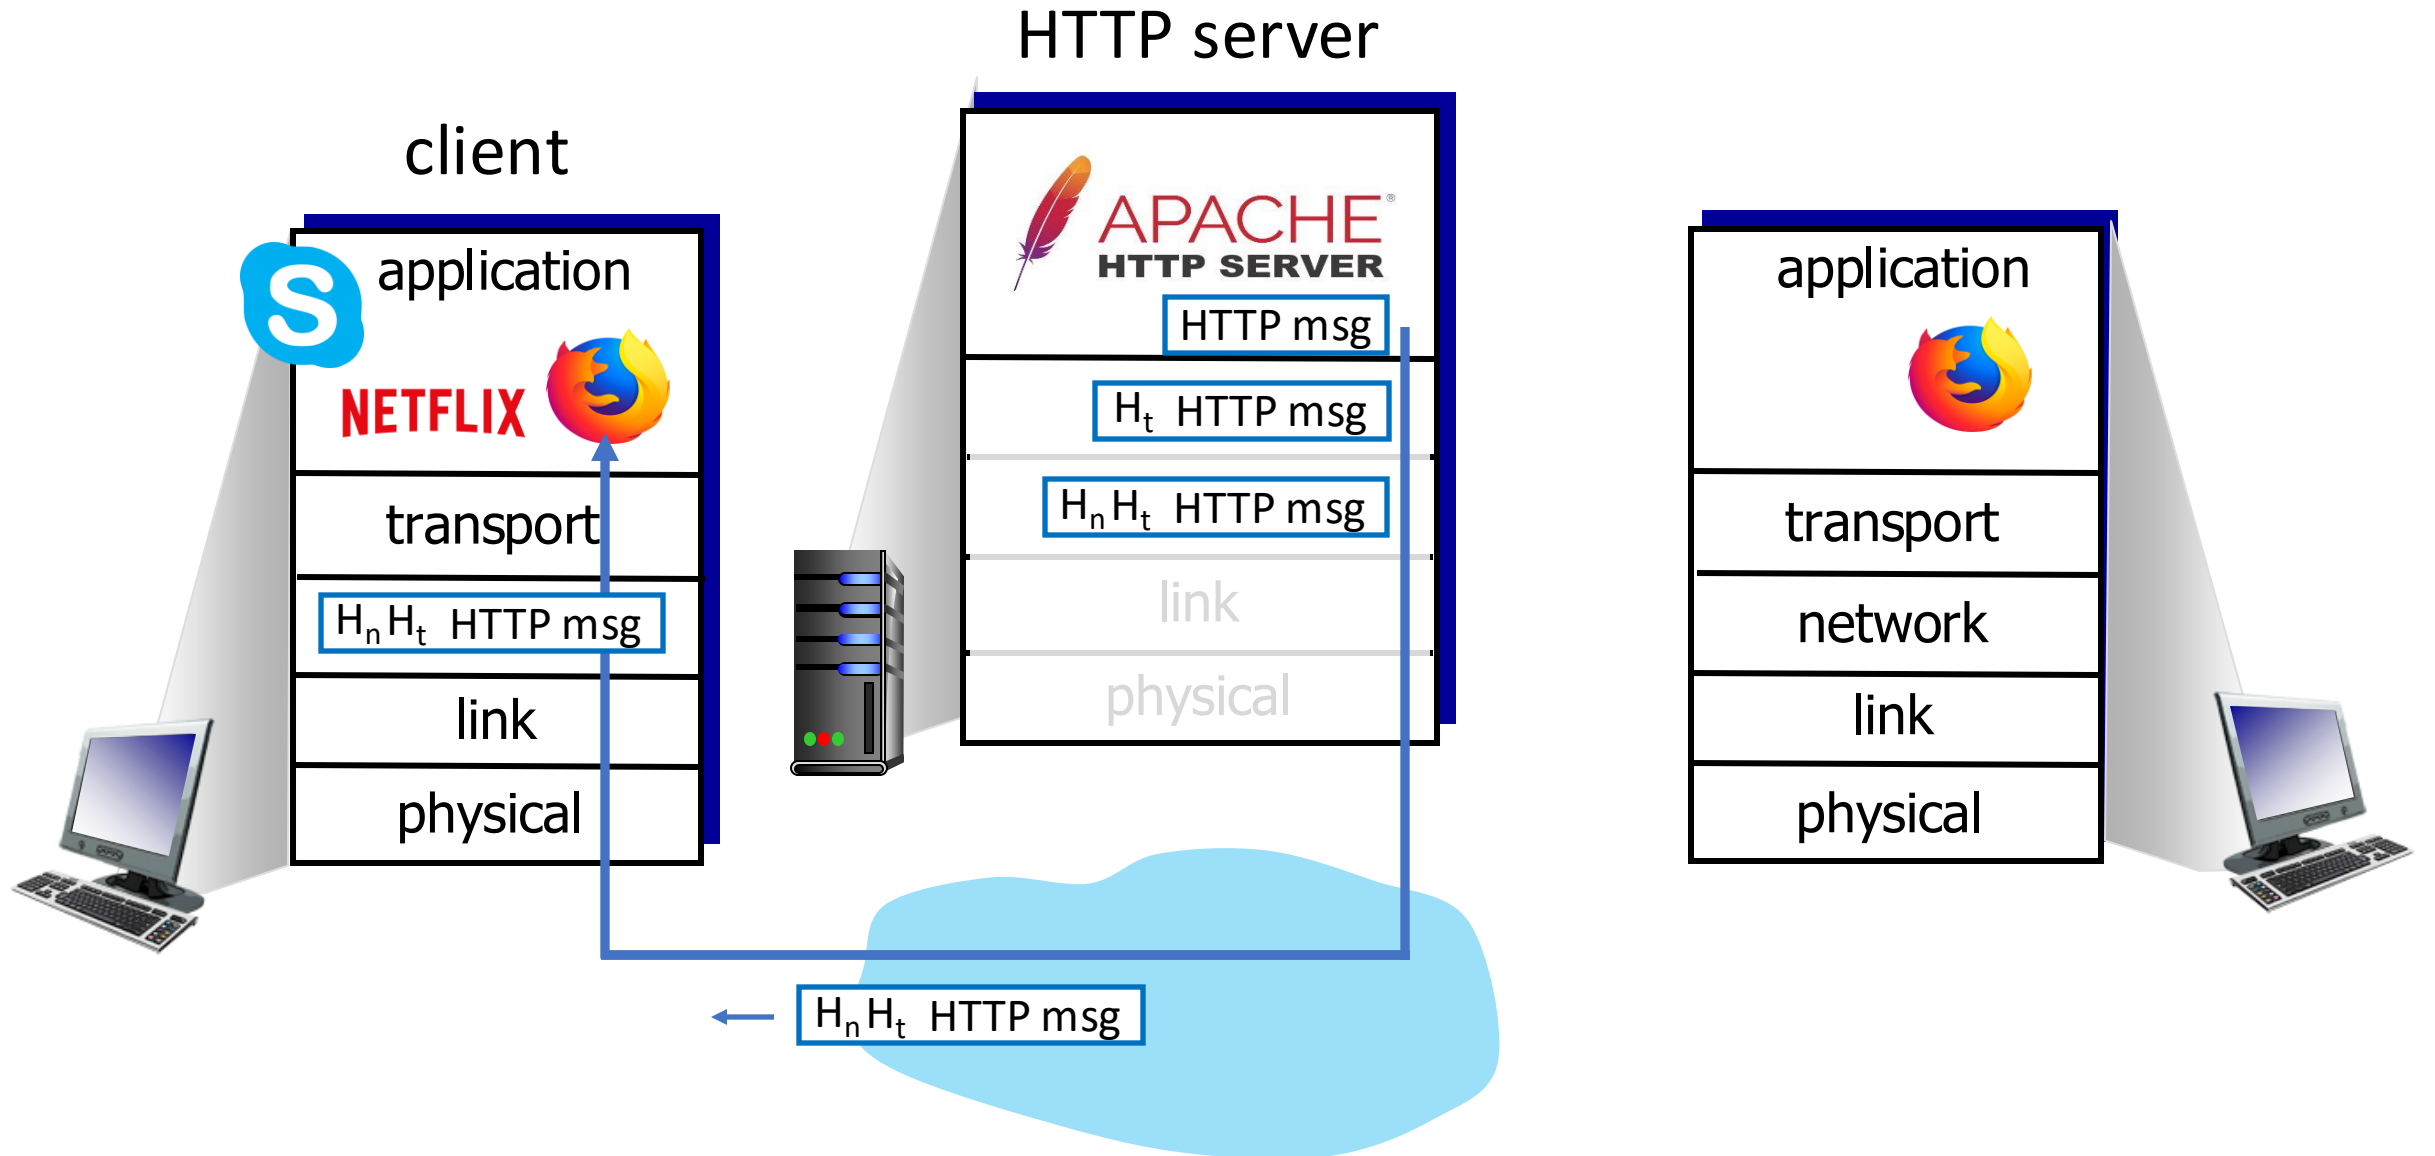
Receiver:

- receives segment from IP
- checks header values
- extracts application-layer message
- **demultiplexes** message up to application via socket




Transport-layer services

- Multiplexing/demultiplexing
- Reliable delivery
- Flow control
- Congestion control






Q: how did transport layer know to deliver message to Firefox browser process rather than Netflix process or Skype process?



How demultiplexing works

- host receives IP datagrams
 - each datagram has source IP address, destination IP address
 - each datagram carries one transport-layer segment
 - each segment has source, destination port number
- host uses *IP addresses & port numbers* to direct segment to appropriate socket



TCP/UDP segment format

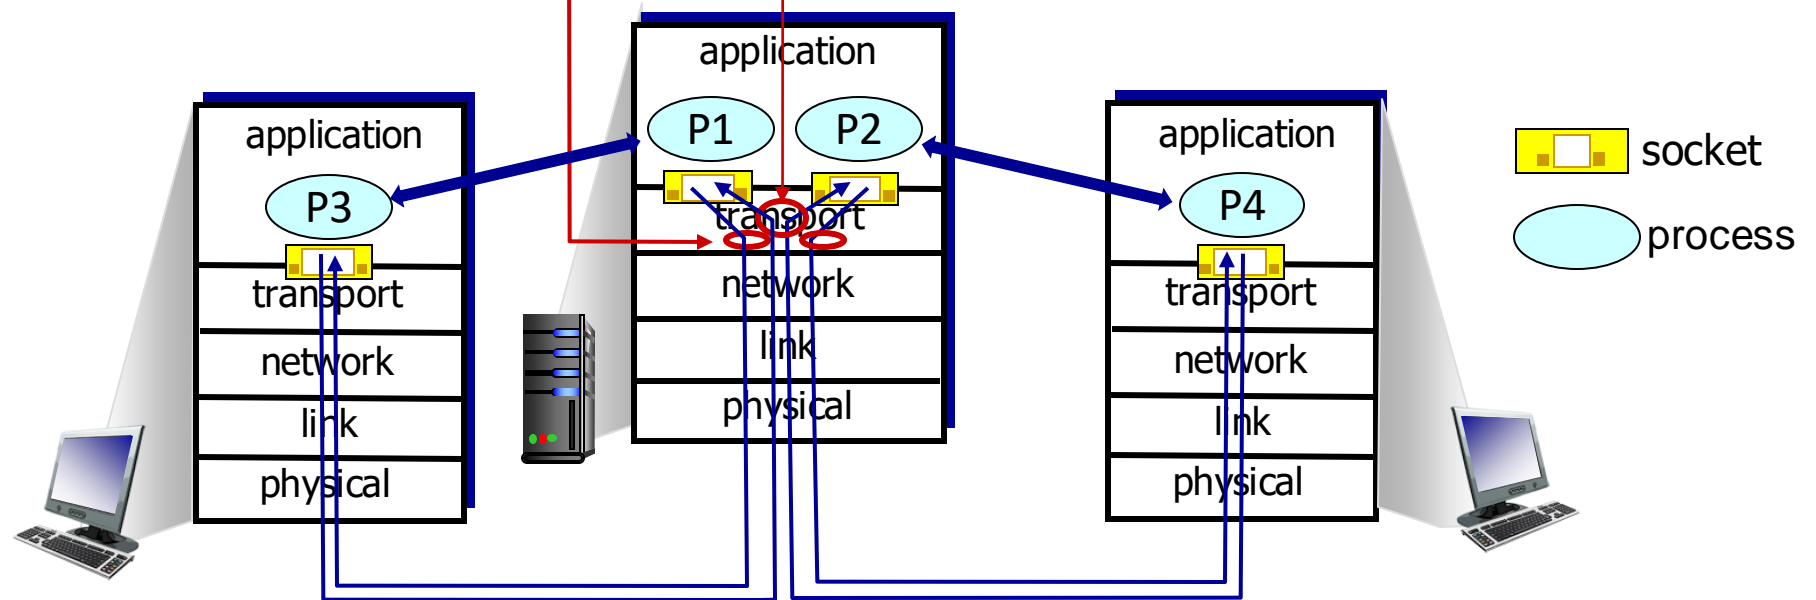
Multiplexing/demultiplexing

multiplexing as sender:

handle data from multiple sockets, add transport header (later used for demultiplexing)

demultiplexing as receiver:

use header info to deliver received segments to correct socket



Multiplexing/Demultiplexing

- Implemented using **logical** ports
- Difference in UDP/TCP multiplexing
- UDP or connection-less: Uses 2-tuple
 - (Dst IP, Dst Port) of the incoming packet aka bind (IP, port)
- TCP or connection-oriented: Uses 4-tuple
 - (Src IP, Src Port, Dst IP, Dst Port)

UDP: User Datagram Protocol

- “no frills,” “bare bones”
Internet transport protocol
- “best effort” service, UDP
segments may be:
 - lost
 - delivered out-of-order to app
- *connectionless*:
 - no handshaking between UDP
sender, receiver
 - each UDP segment handled
independently of others

Why is there a UDP?

- No connection establishment delay. Makes it particularly suitable for applications involving short transactions:
 - DNS
 - SNMP
 - DHCP
- Provides applications with the flexibility to implement only a subset of features
 - E.g., reliability is not needed for video conferencing applications

UDP: User Datagram Protocol [RFC 768]

INTERNET STANDARD

RFC 768

J. Postel

ISI

28 August 1980


User Datagram Protocol

Introduction



This User Datagram Protocol (UDP) is defined to make available a datagram mode of packet-switched computer communication in the environment of an interconnected set of computer networks. This protocol assumes that the Internet Protocol (IP) [1] is used as the underlying protocol.

This protocol provides a procedure for application programs to send messages to other programs with a minimum of protocol mechanism. The protocol is transaction oriented, and delivery and duplicate protection are not guaranteed. Applications requiring ordered reliable delivery of streams of data should use the Transmission Control Protocol (TCP) [2].


Format



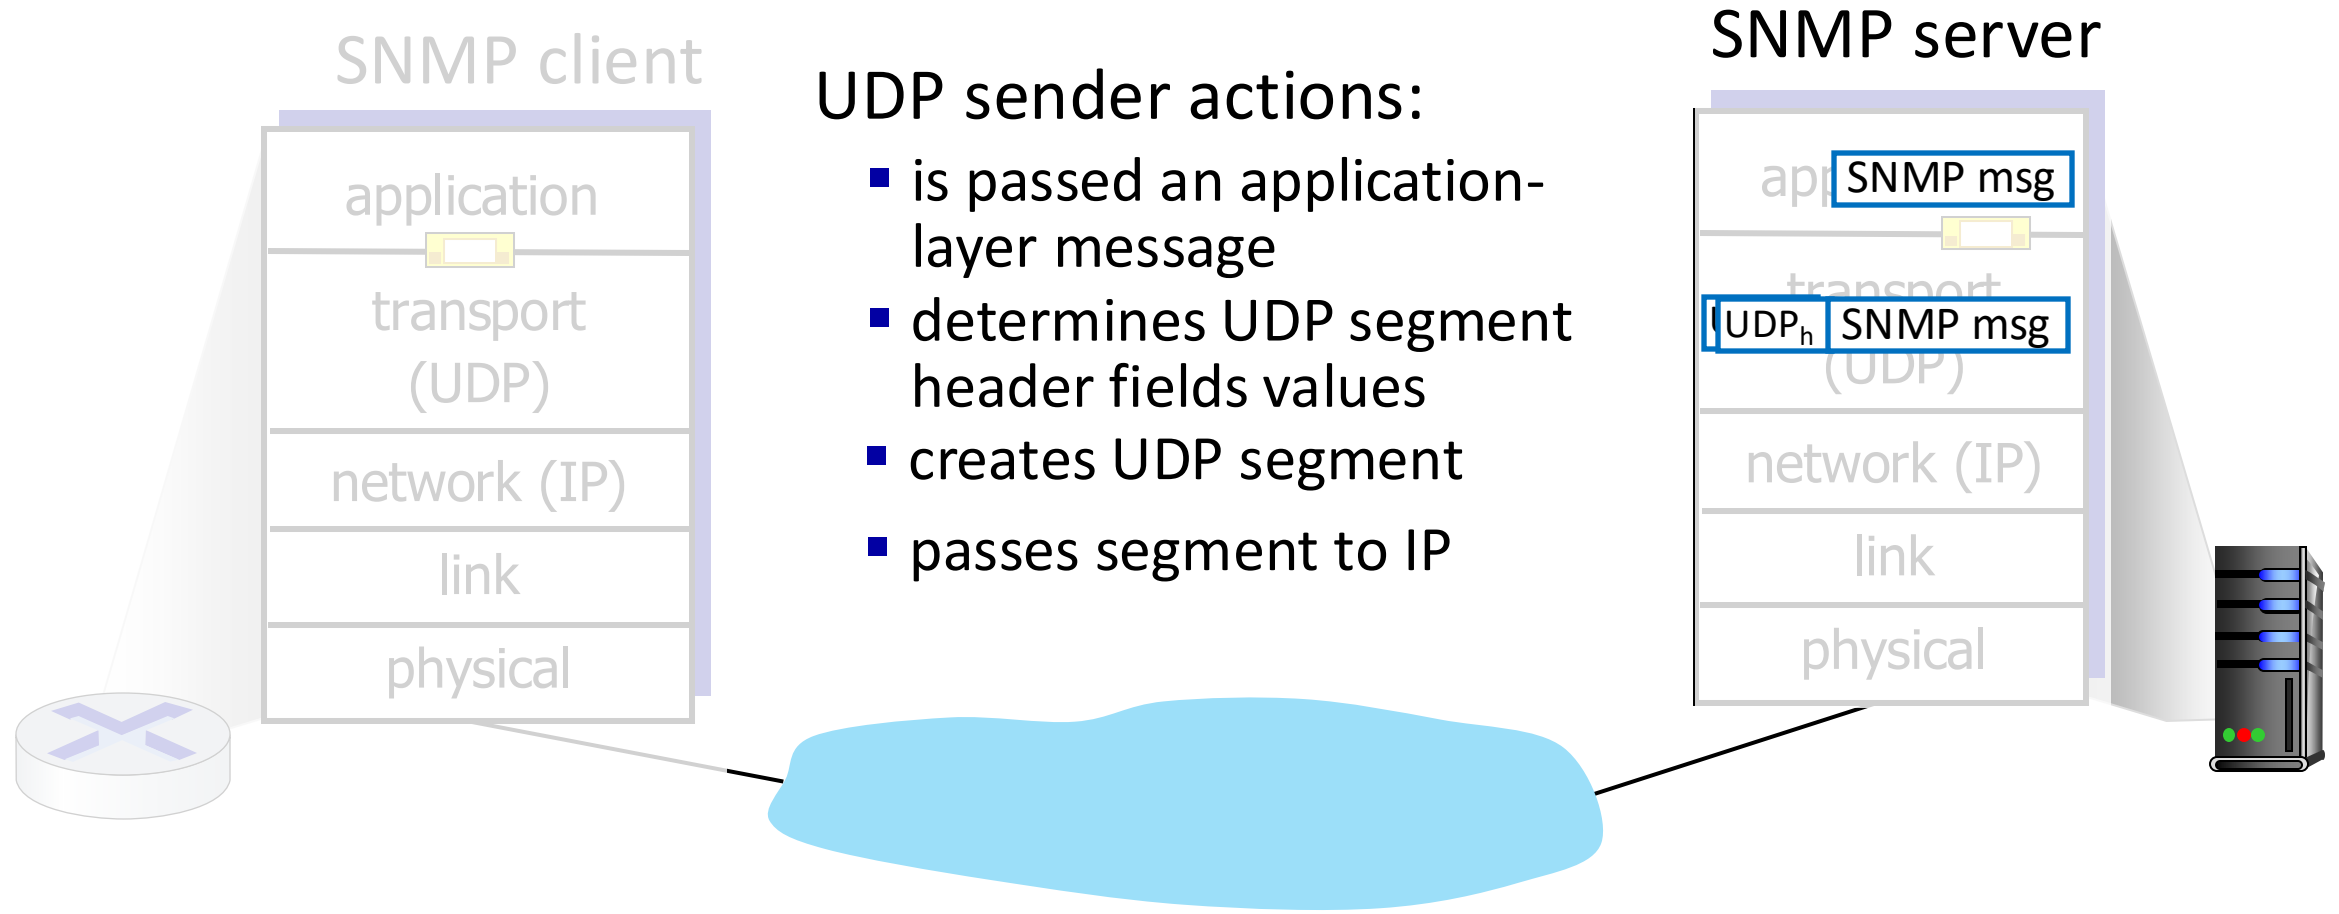
UDP segment header



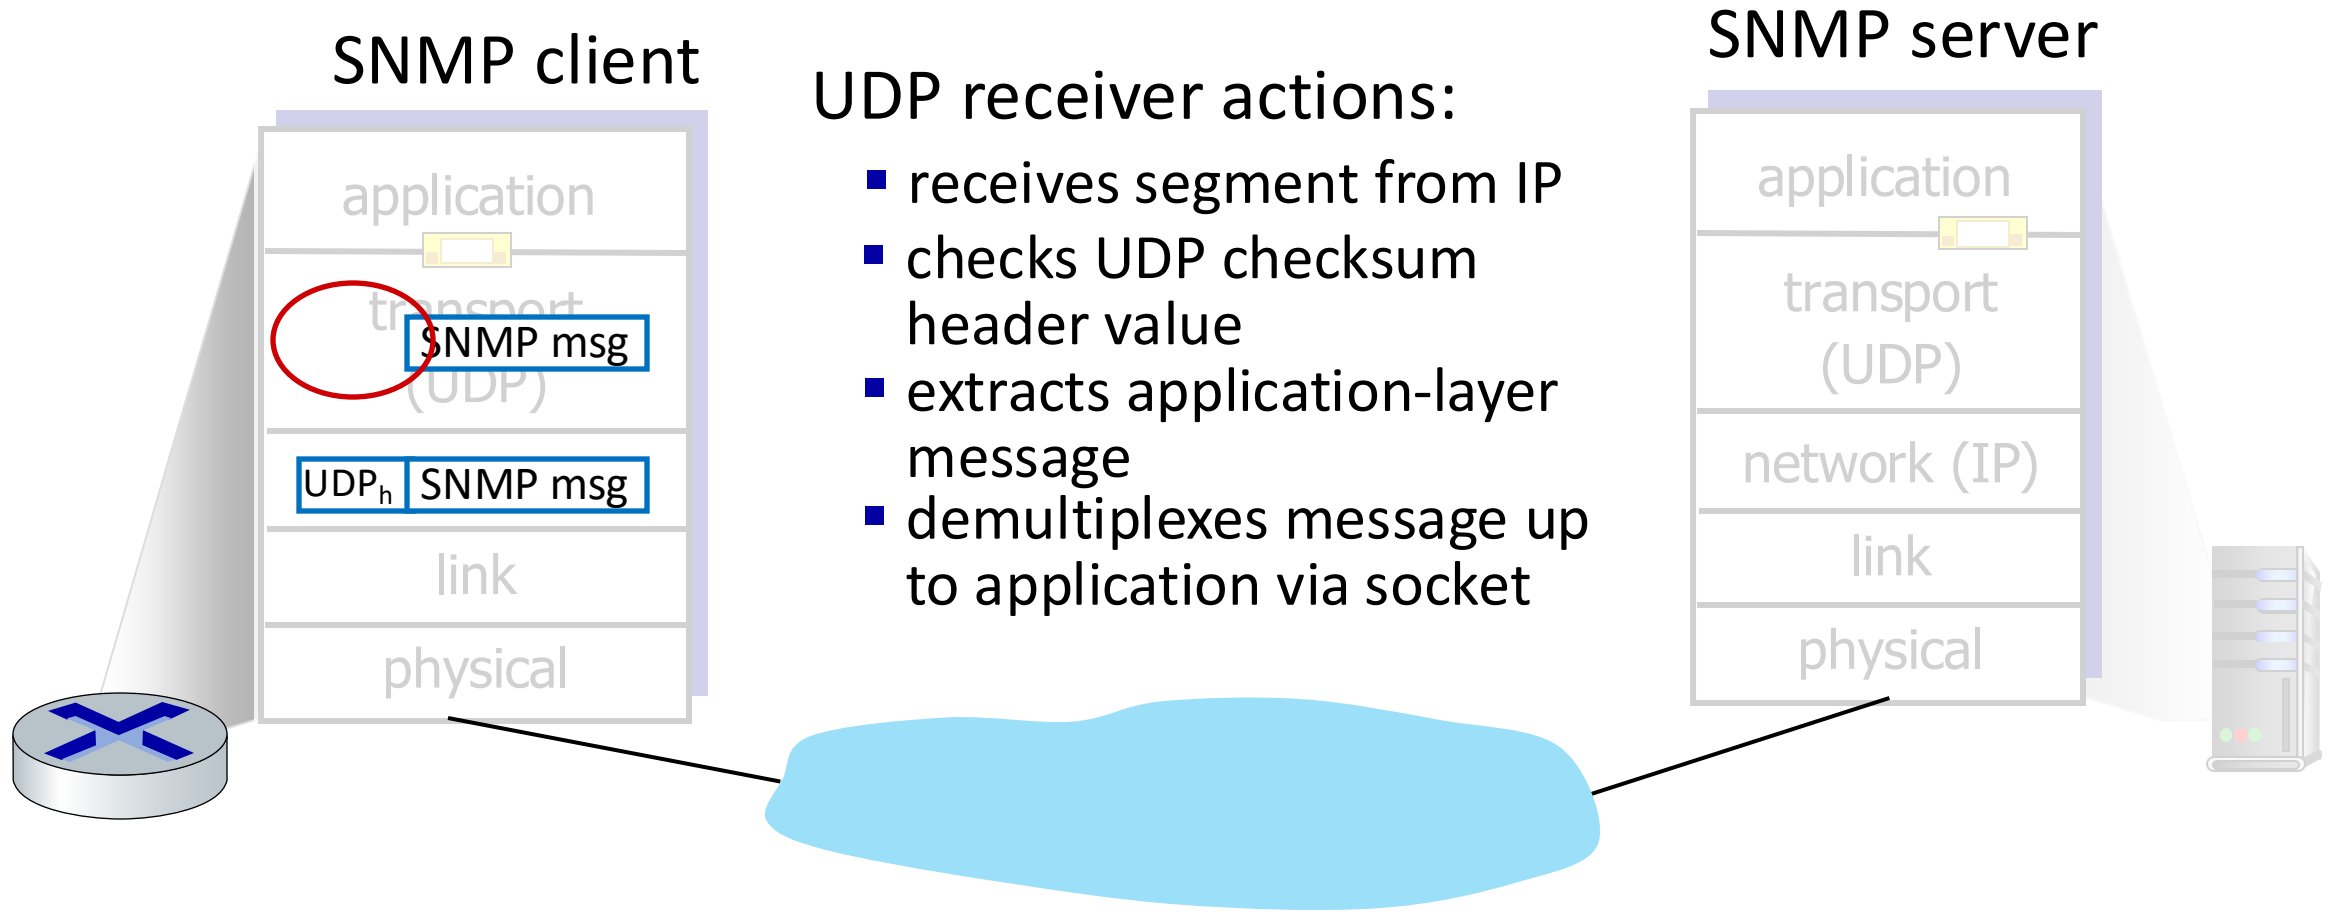
UDP: Transport Layer Actions



UDP: Transport Layer Actions



UDP: Transport Layer Actions



Transport-layer services

- Multiplexing/demultiplexing
- **Reliable delivery**
- Flow control
- Congestion control

Principles of reliable data transfer

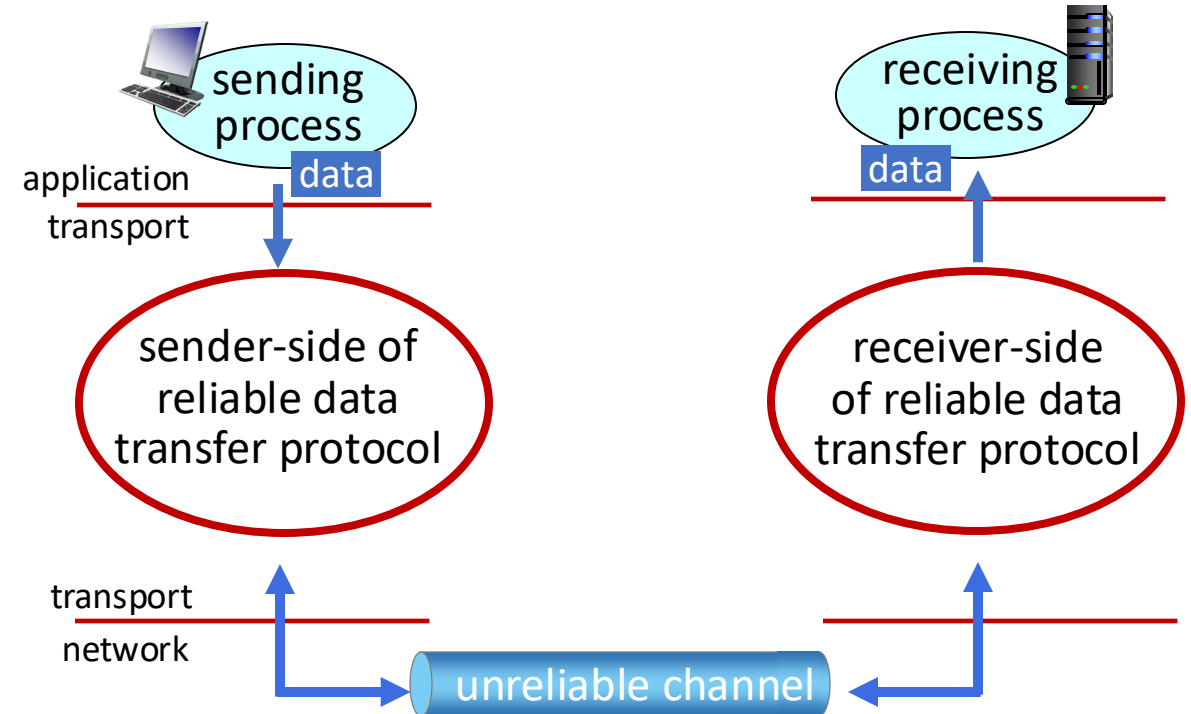



reliable service *abstraction*

Principles of reliable data transfer



reliable service *abstraction*



reliable service *implementation*

Reliability

- Error correction codes
- Automatic Repeat reQuest (ARQ)

Error correction code

- Also known as **Forward Error Correction**


- Using 2D parity

Can detect *and* correct errors
(without retransmission!)

- detect *and correct* single bit errors

- Always useful?

- When cost of retransmissions are high
- When there are frequent bit errors






ARQ Protocol: Stop and Wait

- Transmit one frame, wait for an acknowledgement
 - If no ack and timer expires, resend

Stop and Wait

- Transmit one frame, wait for an acknowledgement
 - If no ack and timer expires, resend
- How to handle duplicate frames?
 - Sequence numbers for duplicate frames
- Any limitation?
 - Under-utilization of link
 - Example, 4 Mbps link, RTT – 10ms, Frame size – 1 KB
 - How to achieve full-link utilization?
 - Bandwidth delay product

Sliding Window Protocol



Summary

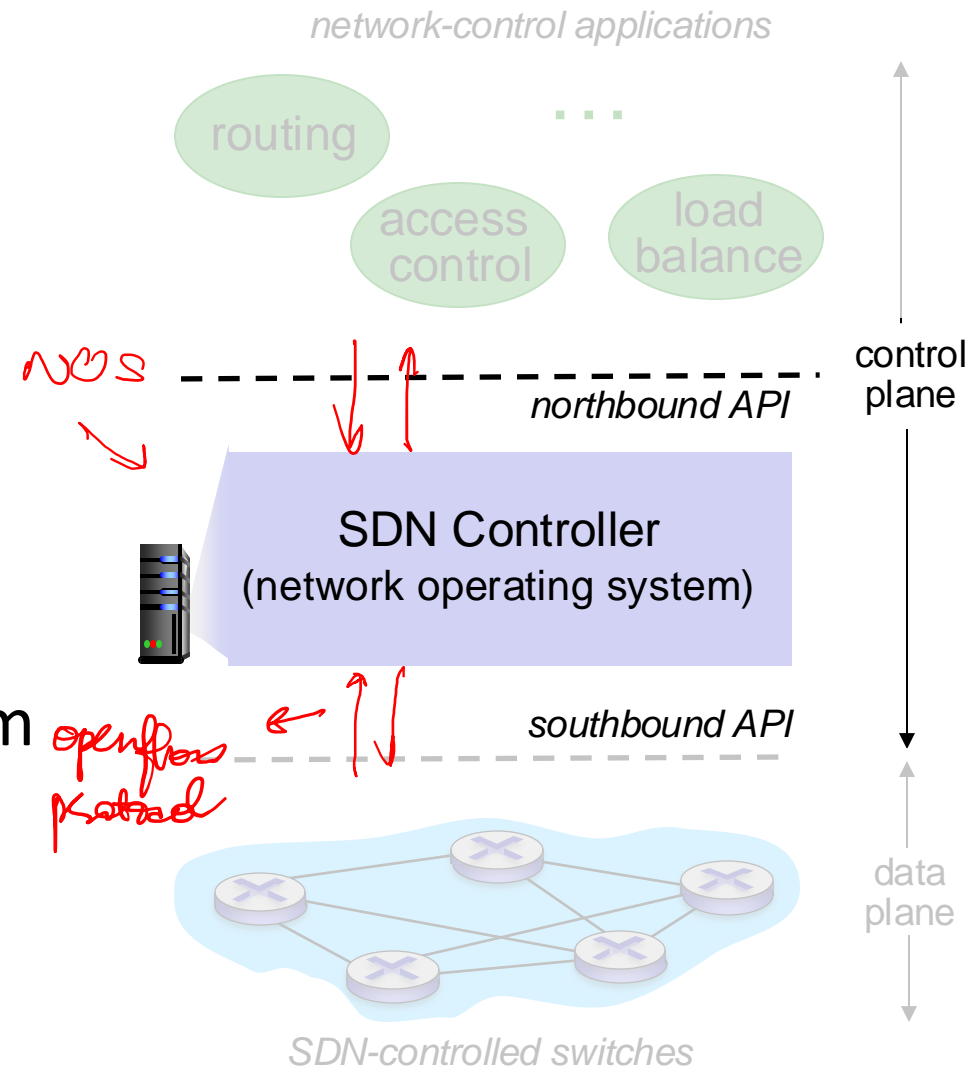
- Transport-layer services
- Multiplexing/demultiplexing
- UDP
- Reliability

Attendance



Recap: SDN Controller


- maintains network state information
- interacts with network control applications “above” via northbound API
- interacts with network switches “below” via southbound API
- implemented as distributed system for performance, scalability, fault-tolerance, robustness




protocol → application

Software defined networking (SDN)

- operators don't "program" switches by creating/sending OpenFlow messages directly.
- Instead use higher-level abstraction at controller
- "brains" of control: implement control functions using lower-level services, API provided by SDN controller
- *unbundled*: can be provided by 3rd party: distinct from routing vendor, or SDN controller




SDN: control/data plane interaction example



- ① S1, experiencing link failure uses OpenFlow port status message to notify controller
- ② SDN controller receives OpenFlow message, updates link status info
- ③ Dijkstra's routing algorithm application has previously registered to be called when ever link status changes. It is called.
- ④ Dijkstra's routing algorithm access network graph info, link state info in controller, computes new routes

SDN: control/data plane interaction example





- ⑤ link state routing app interacts with flow-table-computation component in SDN controller, which computes new flow tables needed
- ⑥ controller uses OpenFlow to install new tables in switches that need updating

SDN: Key Challenges

■ Hardening the control plane

- Scalability
 - Reliability
 - Consistency
 - Security
- Real-time → logically centralized
- Distributed systems



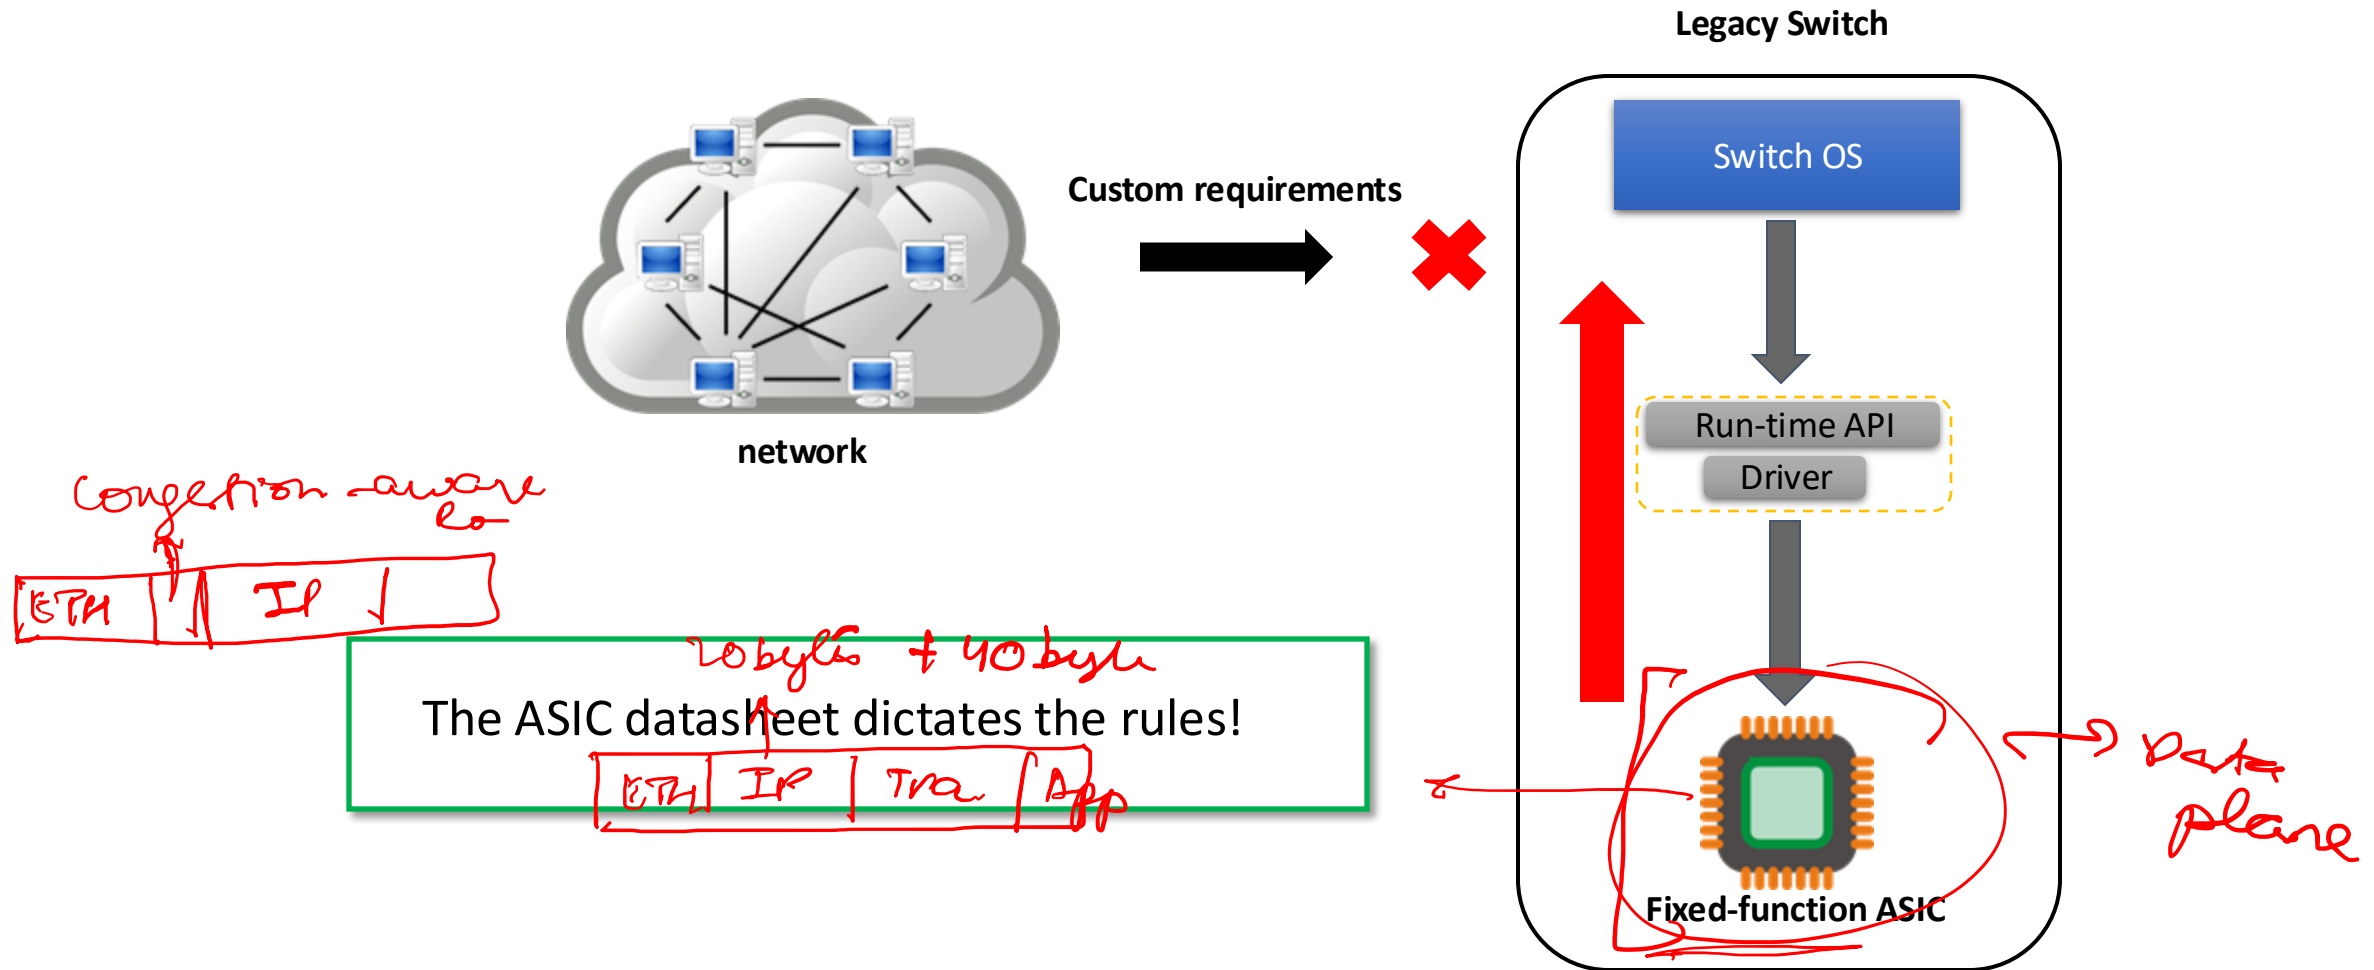
■ Internet-scaling: beyond a single AS (?)


What else is programmable in the network?

- Programmable data plane }
- Network function virtualization (NFV)

↓
making middleboxes programmable

Status Quo: Bottom-up design








How to make Data Plane Programmable?


■ Move the data plane to software

- Generally, too slow for data plane functions
- Any optimization techniques?
 - Parallelism across multiple servers and cores
 - Optimizations in NUMA
 - Fast I/O ...





■ What about programmable hardware?

- FPGA: but costly, power-hungry, slower



How to make Data Plane Programmable?

- Move the data plane to software
 - Generally, too slow for data plane functions
 - Any optimization techniques?
 - Parallelism across multiple servers and cores
 - Optimizations in NUMA
 - Fast I/O ...
- What about programmable hardware?
 - FPGA: **but costly, power-hungry, slower**

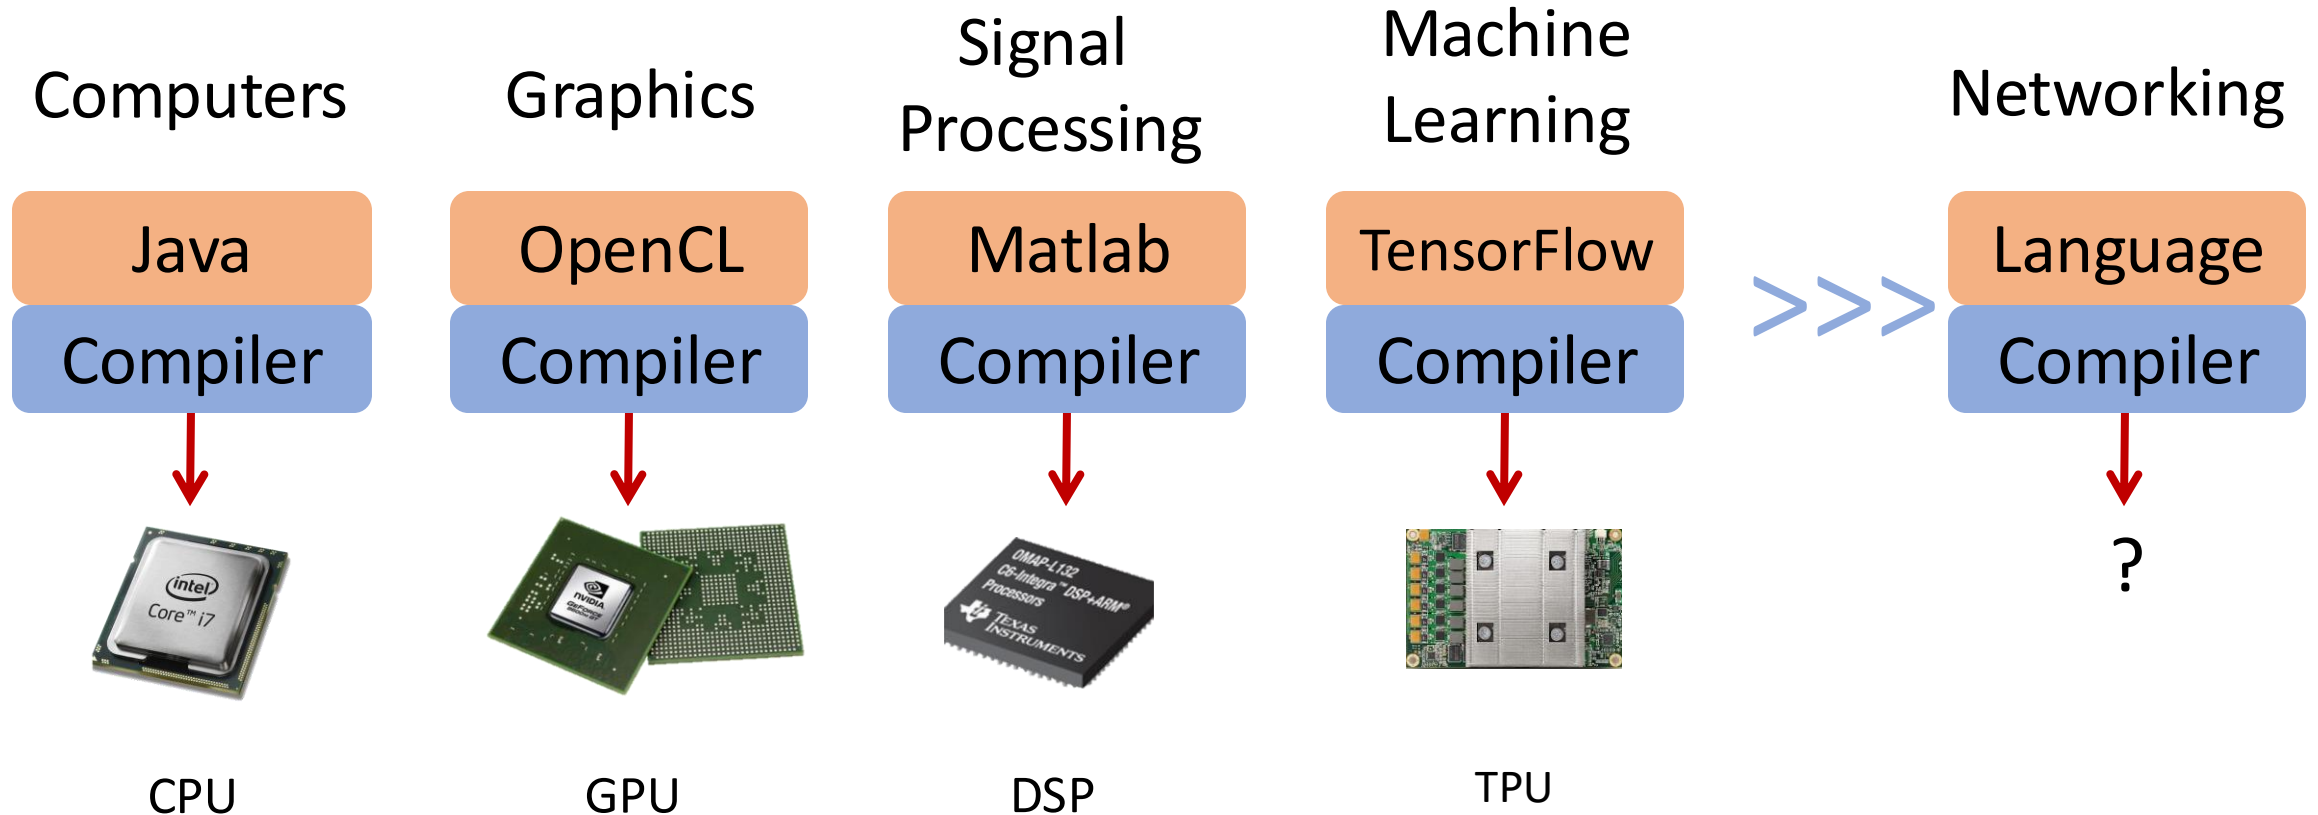


What about Programmable Hardware?

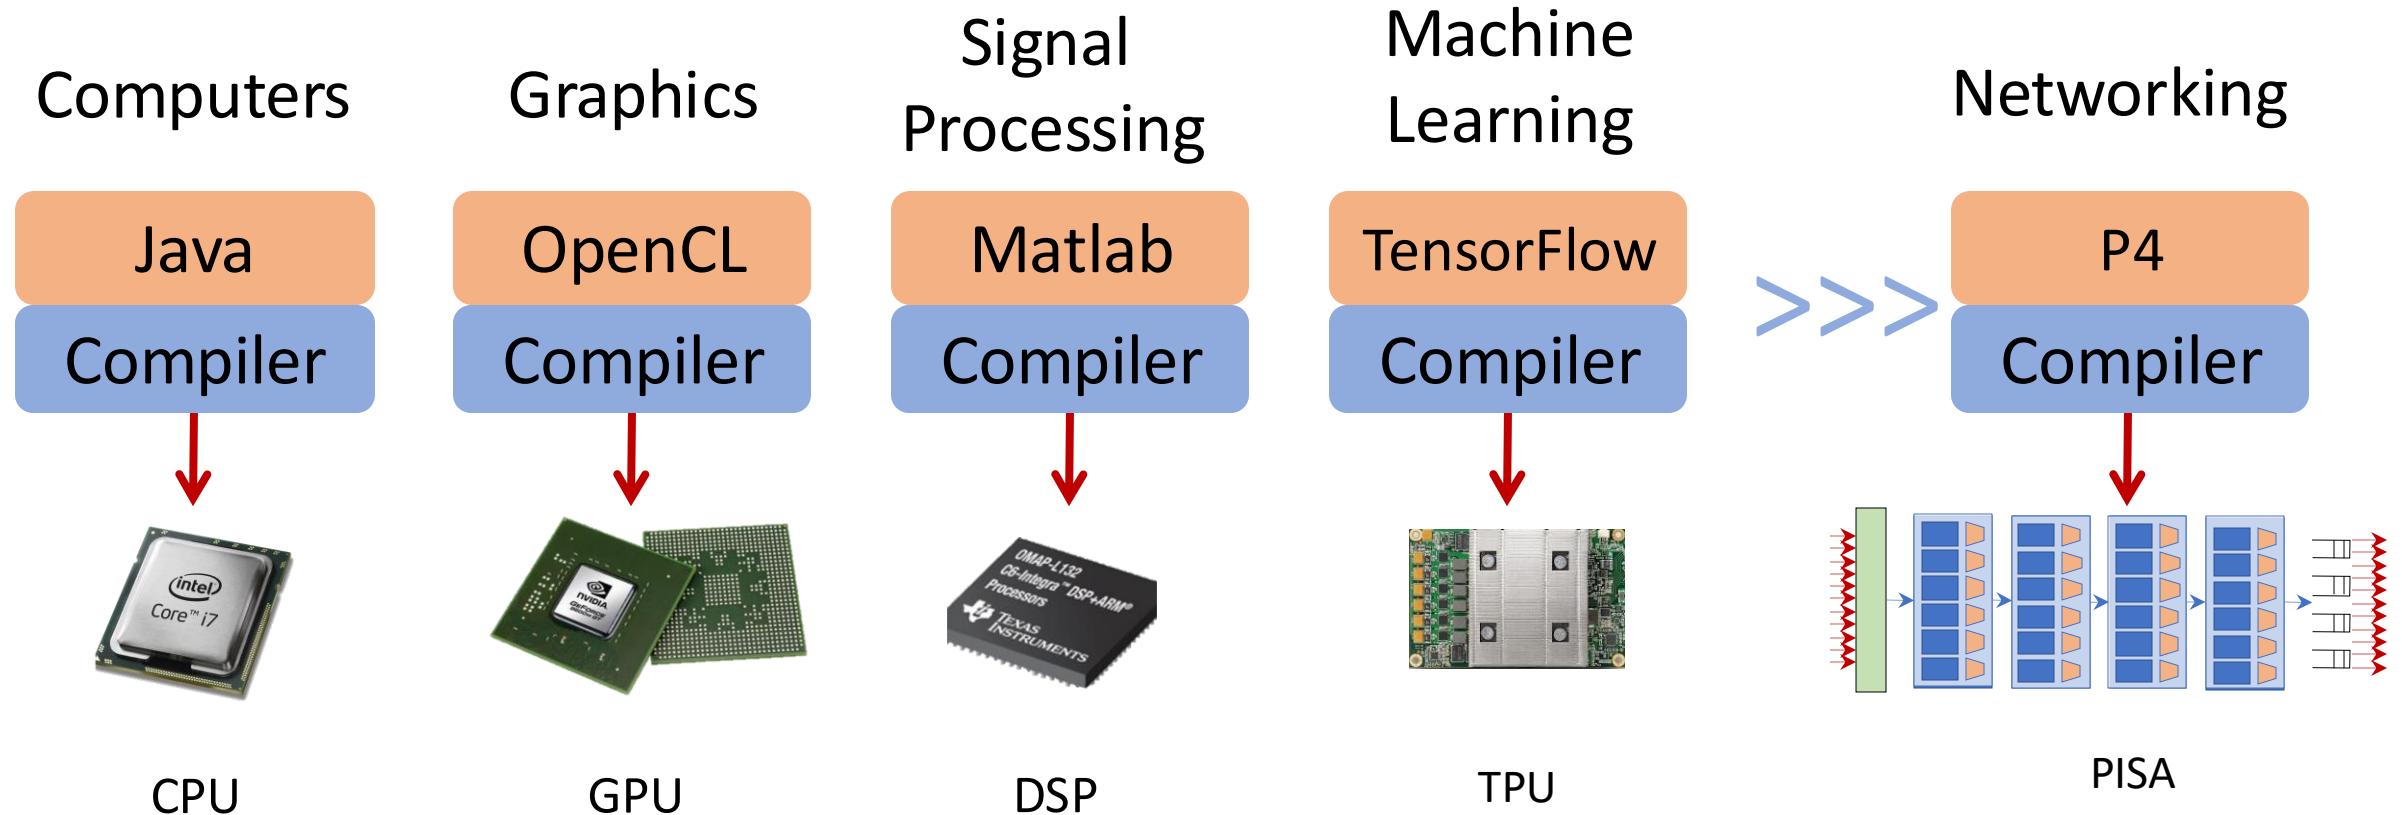
“Programmable switches are 10-100x slower than fixed-function switches. They are more expensive and consume more power.”

Conventional wisdom in networking

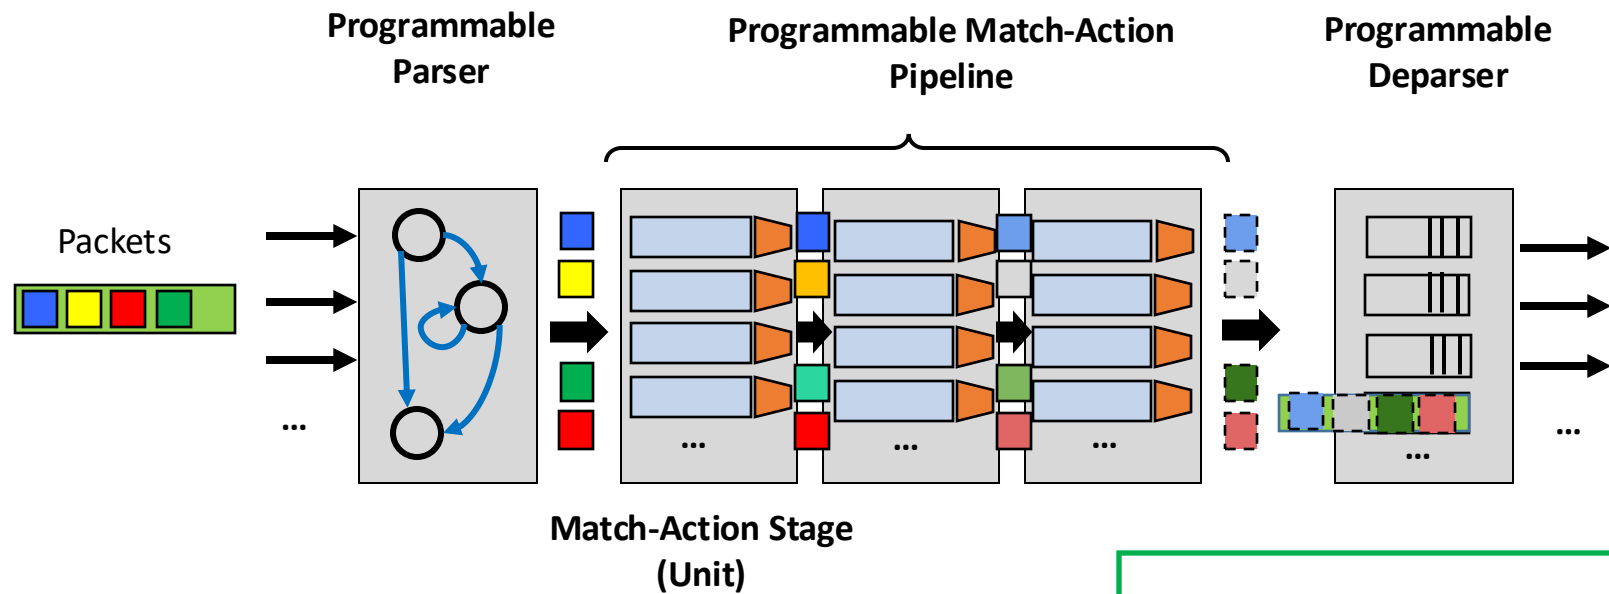
Domain Specific Processors



Domain Specific Processors



Protocol Independent Switch Architecture



P4 describes PISA's behavior

Reducing complexity

switch.p4

Switch OS

IPv4 and IPv6 routing

- Unicast Routing
 - Routed Ports & SVI
 - VRF
- Unicast RPF
 - Strict and Loose

~~Multicast~~

- ~~- PIM-SM/DM & PIM-Bidir~~

Ethernet switching

- ~~- VLAN Flooding~~
- MAC Learning & Aging
- STP state
- ~~- VLAN Translation~~

Load balancing

- ~~- LAG~~
- ECMP & WCMP
- Resilient Hashing
- ~~- Flowlet Switching~~

Fast Failover

- LAG & ECMP

Tunneling

- IPv4 and IPv6 Routing & Switching
 - ~~- IP in IP (Gin4, 4in4)~~
 - VXLAN, NVGRE, GENEVE & GRE
 - ~~- Segment Routing, ILA~~

~~MPLS~~

- ~~- LER and LSR~~
- ~~- IPv4/v6 routing (L3VPN)~~
- ~~- L2 switching (EoMPLS, VPLS)~~
- ~~- MPLS over UDP/GRE~~

ACL

- MAC ACL, IPv4/v6 ACL, RACL
- ~~- QoS ACL, System ACL, PBR~~
- Port Range lookups in ACLs

QoS

- QoS Classification & marking
- ~~- Drop profiles/WRED~~
- ~~- RoCE v2 & FCoE~~
- CoPP (Control plane policing)

~~NAT and L4 Load Balancing~~

Security Features

- ~~- Storm Control, IP Source Guard~~

Monitoring & Telemetry

- ~~- Ingress Mirroring and Egress Mirroring~~
- Negative Mirroring
- ~~- Sflow~~
- INT

Counters

- Route Table Entry Counters
- ~~- VLAN/Bridge Domain Counters~~
- Port/Interface Counters

Protocol Offload

- BFD, OAM

Multi-chip Fabric Support


- ~~- Forwarding, QoS~~

What else is programmable in the network?

- Programmable data plane
- **Network function virtualization**

Middleboxes

Data delivery is not the only required functionality.




Elements in the network path for security, performance enhancements etc.


One-third of all network devices in enterprises are middleboxes!
Sherry et al., SIGCOMM'12

Evolution of Middleboxes


Dedicated hardware



*Need for
flexibility*


A black arrow pointing from the dedicated hardware side to the software side.

Software




From Hardware Middleboxes..

Data delivery is not the only required functionality.



To Software Network Functions (NF)



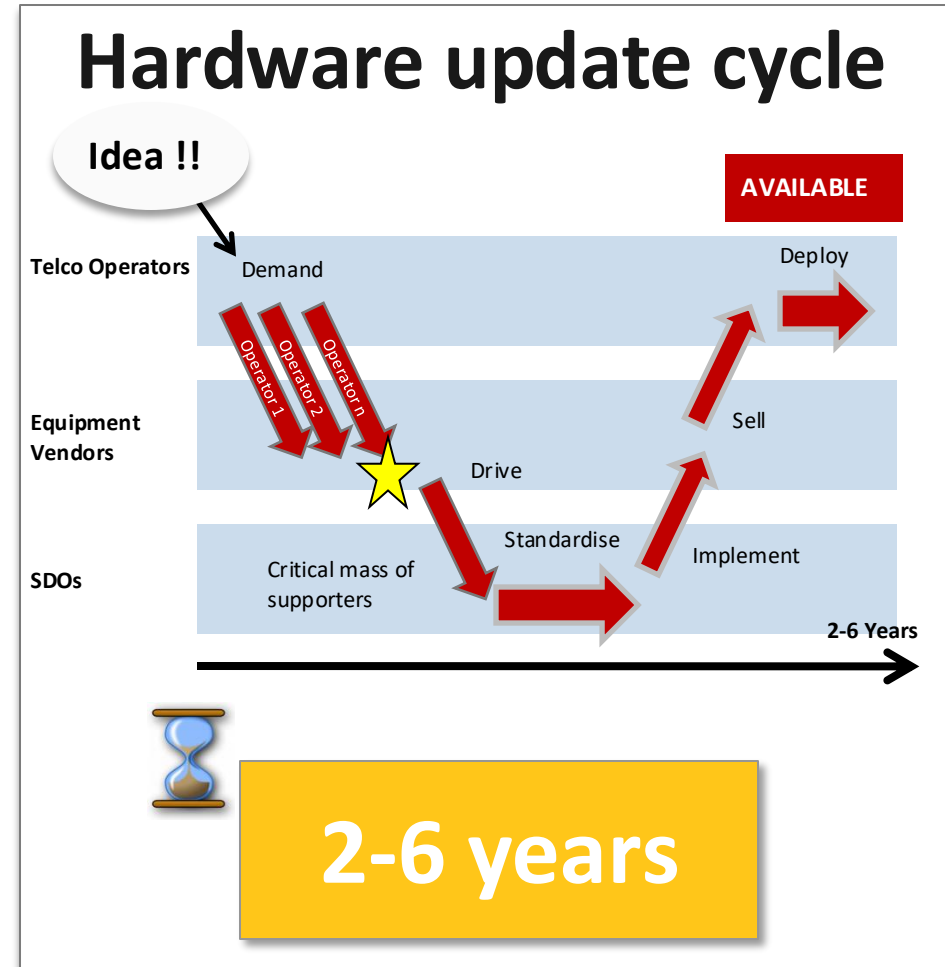
Primarily deployed in a VM
(Network Function Virtualization or NFV)

Functional Elements, not Middleboxes

- **WAN Optimizer** = Caching + Deduplication + Compression + Encryption + Forward Error Correction + Rate Limiter
- **Application Firewall** = IP Defragmenter + Application Detection Engine + Logger + Blocker
- **IDS** = IP Defragmenter + Preprocessing + Misuse Detection Engine + Logger


Why NFV?

- Softwarization leads to faster innovation



Why NFV?

- Softwarization leads to faster innovation
- Ease of deployment, configuration, and management
- Consolidation: Reduce number of hardware boxes in the network



Being adopted by both carriers and cloud providers

NFV Challenges

- Virtual network function management
 - Where and how to install network functions?
- Unpredictable (low) Performance
 - How to mitigate the overheads of virtualization?
- Fault Tolerance
 - How to handle recovery in case of faults?

Summary

- Increased programmability in the networks
 - Greater flexibility → Faster innovation
- Programmable control plane
 - SDN
- Programmable data plane
 - SDN-2 / P4
- Software middleboxes implemented in VMs
 - Network function virtualization (NFV)

Attendance

