

Computer Networks

COL 334/672

Application Layer: DNS and P2P

Tarun Mangla

Slides adapted from KR

Sem 1, 2024-25


Recap

- Peer-to-Peer (P2P) networks for content distribution:

- • Scale better as they can make use of client uplink
- Particularly popular in the early 2000s
Napster, Gnutella, BitTorrent

- Two interesting questions:

- ① → • How to find content?
- How to download content?




Finding a File: Approaches

✓NAPSTER (architecture)

- Approach #1: use indexing a centralized server
 - The centralized server contains information about nodes and the files
 - A new node communicates with the centralized server for file search
 - Cons:
 - Single point of failure
 - Accountable
- Approach #2:
 - Node broadcasts query to its neighbors which in turn broadcast it to their neighbors
 - Use TTLs to avoid indefinite broadcast messages
 - Cons: high overhead

Can we do better?




Finding a file in a P2P network

hash table: File \rightarrow which peer

- ① Distributed manner
- ②

- **Intuition:** Some indexing is useful for a faster lookup. What kind?
- **Challenge:** But can't have a centralized hash table
- **Solution:** Use a distributed hash table (DHT)




PASTRY

Idea:

- Map the objects and the nodes to a common virtual space
- Store the object information in a node that is closest to it in the abstract space

How do we search for the closest node?



PASTRY

Idea:

- Map the objects and the nodes to a common virtual space : If addr space
- Store the object information in a node that is closest to it in the abstract space


① what if the peer disconnects

$H(P_i)$ →

Peers
may not
be unfor

How do we search for the closest node?

$2^3 - 1$
 $H(P)$ 0, 2, 6
 $H(F)$ 2, 3, 5
 $F_1 : 6$
 $F_2 : 2$
 $F_3 : 0$



()

PASTRY: Searching Closest Node

Idea: To search for a file f , route query messages closer to $H(f)$ in the virtual space until you find the node containing information about f


Challenge: How do we ensure that we can always go to a closer node?

$$L = 2$$



Solution: Each node should store L nodes ($L/2$ successors, $L/2$ predecessors) and $\log(N)$ nodes distributed randomly in the virtual space (Randomized algorithm)

$\log(N)$ chord

$$H(F_1) = 2 \quad H(F_2) = 3 \quad H(F_3) = 5$$



Dynamically change



$$2^m + 1$$

PASTRY: Searching Closest Node

Idea: To search for a file f , route query messages closer to $H(f)$ in the virtual space until you find the node containing information about f


Challenge: How do we ensure that we can always go to a closer node?

$L = 2$



Solution: Each node should store L nodes ($L/2$ successors, $L/2$ predecessors) and $\log(N)$ nodes distributed randomly in the virtual space (Randomized algorithm)

$\log(N)$ chord

$H(F_1) = 2$ $H(F_2) = 3$ $H(F_3) = 5$



Dynamically change



$2^m + 1$

Distributed Hash Table

- You should think about the following:
 - How the neighbor are maintained in the first place
- Various optimizations exist for DHTs
- Used in other domains such as distributed file system, web caching etc.

Next question: How to download content?

P2P file distribution: BitTorrent

- ① Rarest one first
- ② In-order (Application consider)
- ③ Randomly down

File chunks

①. What's content to download from whom?

②. who to send the data?

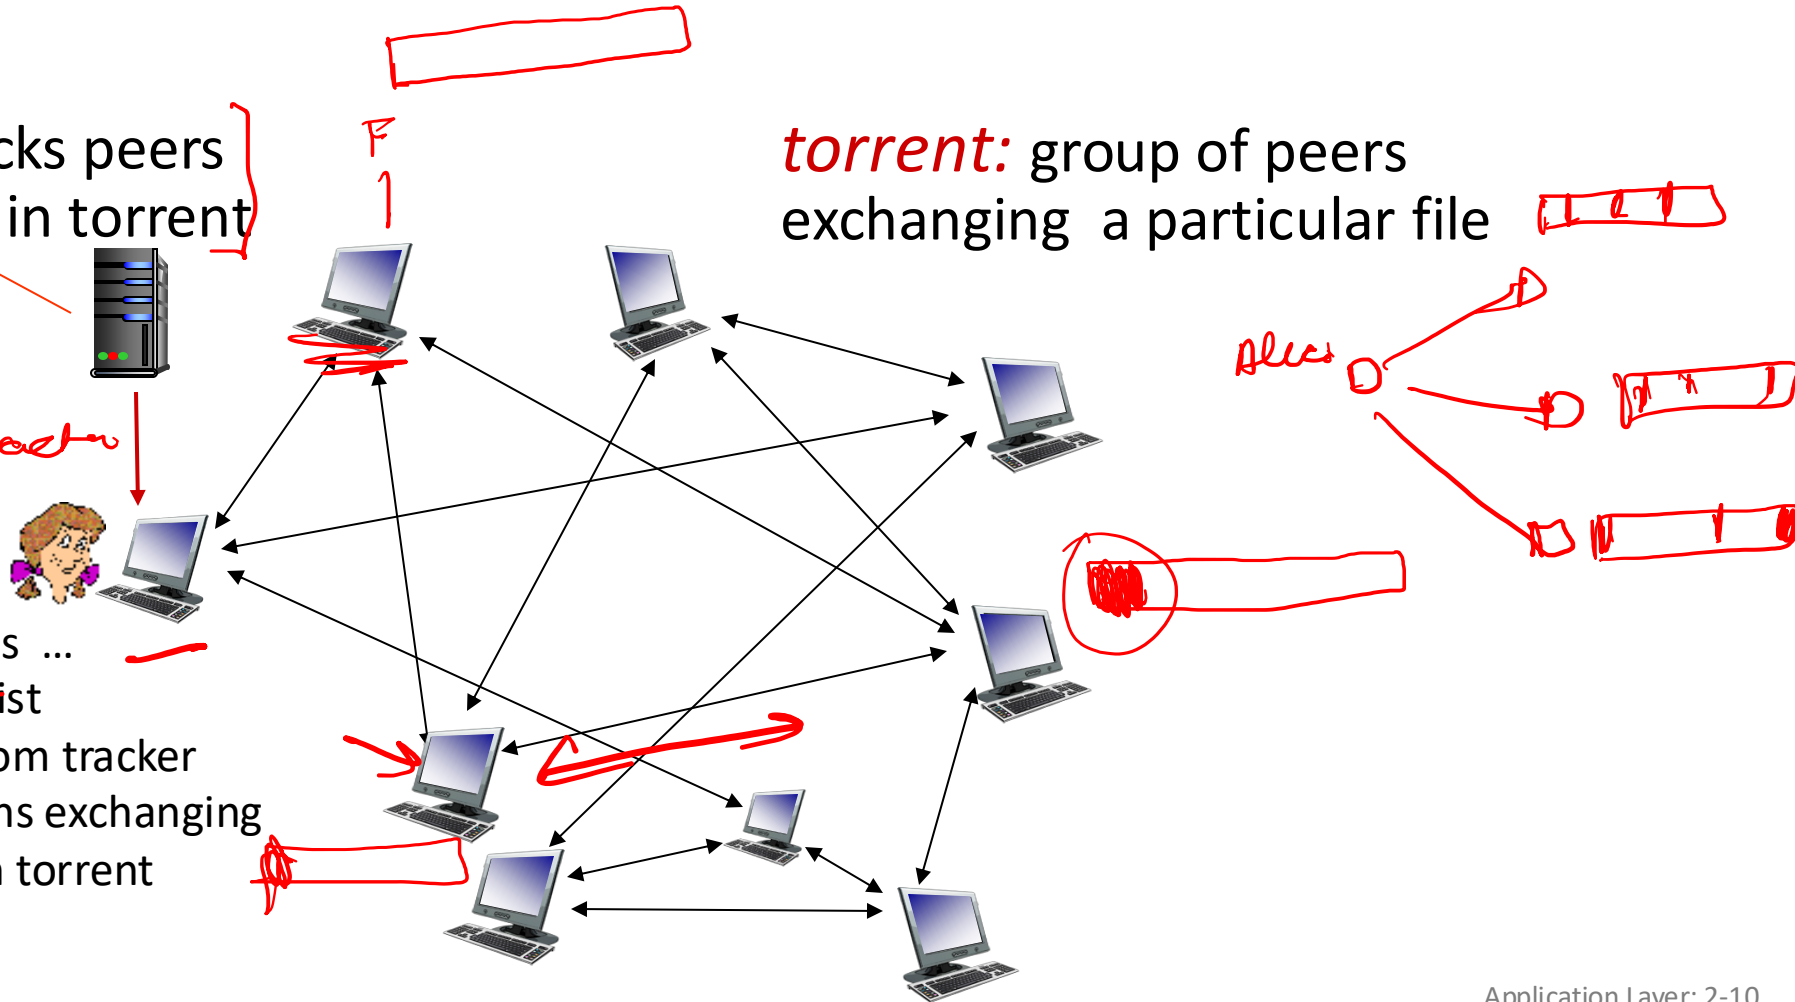
[tracker: tracks peers participating in torrent]

① centralized tracker

② Decentralized tracker (DHTs)

torrent: group of peers exchanging a particular file

Alice arrives ...
... obtains list of peers from tracker ... and begins exchanging file peers in torrent



BitTorrent: requesting, sending file chunks



Which chunks to request?

- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first


Sending chunks: whom to send chunks?

Exploration
and exploitation
W

- Uses tit for tat
- sends chunks to those four peers currently sending chunks *at highest rate*
 - other peers are choked by Alice (do not receive chunks from her)
 - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks

BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



Need for alternate “faster” content distribution mechanism


Content Delivery Networks

- P2P would not work for latency-sensitive applications (e.g., web)
- Need an alternate mechanism that scales well (number of users and geography)
- **Question:** how to scale client-server paradigm?

Use geographically distributed servers

- Too expensive to do that for every content provider!

Use Content Distribution Networks



Content distribution networks (CDNs)

Akamai → MIT

- CDN: geographically distribute collection of *server surrogates*
- Servers can be leased by many customers
- Popular CDNs: Limelight, Akamai, Level3
- Two kinds of server placement policies:

- *enter deep*: push CDN servers deep into many access networks

- close to users
- Akamai: 240,000 servers deployed in > 120 countries (2015)

- *bring home*: smaller number (10's) of larger clusters in POPs near access nets

- used by Limelight

Twitter - 1 JSRs

